

# Minix File System – File Layout

- The Minix FS is a logical, self-contained entity with i-nodes, directories and data blocks. Example of a 360K layout with 127 i-nodes and 1K block size. Figure 5.30.




Figure 5.30: Minix FS Layout

- **Boot block.** When a computer is turned on, the hardware reads the boot block into memory and jump to it.

# Super Block

- **super-block** contains information about file system layout. Figure 5.31.

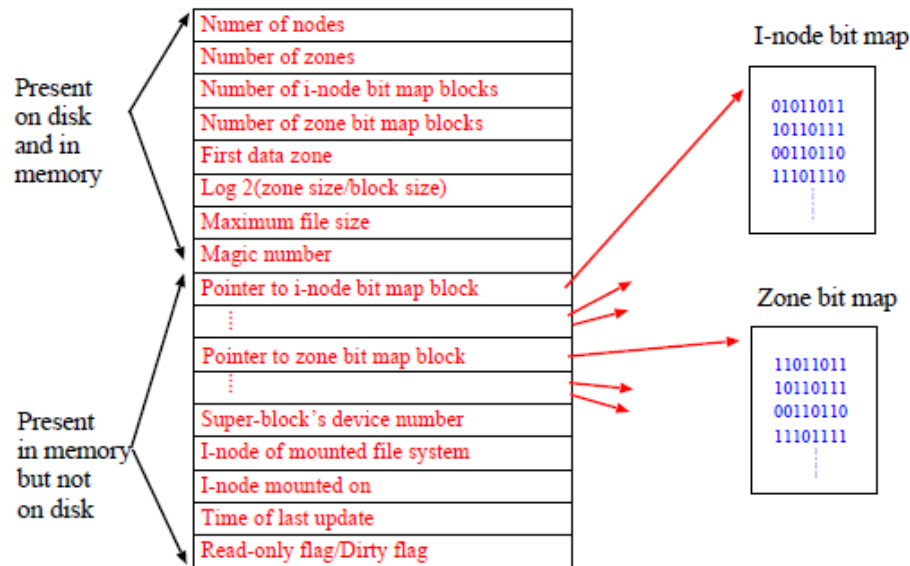


Figure 5.31: The MINIX super-block

# Super Block

- Given block size and number of i-nodes. It is easy to compute the size of i-node bit map and number of blocks for i-nodes. For example, with 1K block size, each block of bit maps can indicate status of 8191 i-nodes. If i-nodes are of 32-bytes, each block holds 32 i-nodes, therefore, 127 i-nodes need 4 disk blocks.

# Super Block

- disk storage is allocated in units (zones) of 1,2,4,8 or in general,  $2^n$  blocks.
- mapping from zone to block or vice versa can be done by algorithm. For example, with 8 blocks per zone, to find zone containing block 128, just shift 128 right by 3 bits.
- when Minix is booted, the super-block for the root device is read into a table called super-blocks table. New parameters are then derived and stored in this table like whether it has been mounted read-only, modified status field.
- To enforce a known structure, the utility program *mkfs* is provided to build a file system.

# Bitmap

- i-nodes and zone bit maps keep track of status.
- upon boot up, super-block and bit maps for the root device are loaded into memory.
- to remove a file, based on the i-node number, we will know which pointer in the super-block to go to and access the i-node bit map and then clear the corresponding bit.
- to create a file, sequentially search through the bit map blocks until it finds a free i-node. This i-node is then allocated for the new file. If no available i-node, return 0.
- Rational for zone is to improve performance so that disk blocks that belong to the same file are located on the same cylinder as much as possible.

# Inode

- in Minix, i-node is a 32-bytes structure. Figure 5.32.

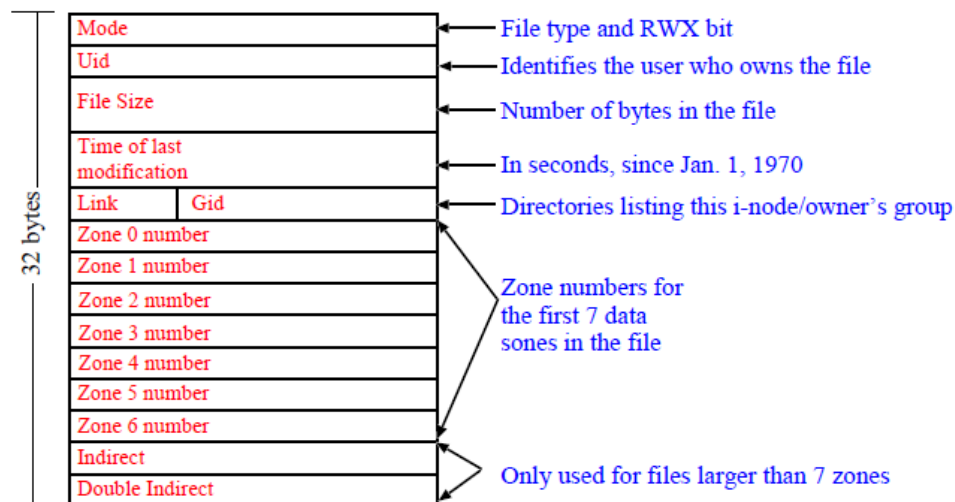


Figure 5.32: The MINIX i-node

- *mode* indicates the type of file (regular, directory, block special, character special or pipe), protection bits.

# Inode

- It is different from traditional unix 1) number of *time* field, 2) link and gid sizes, 3) fewer disk block pointers.
- When a file is opened, its i-node is located and brought into the *inode table* in memory. It remains there until the file is closed.
- the inode table has several parameters which is not in disk, 1) i-node's device number so system knows where to write back the i-node, 2) a counter for each i-node to indicate the number of process opening the file. When the counter is zero, the i-node is removed from the table and written to disk if it has been modified.

# Directories and Path

- as mentioned before, when the system wants to open a file, it is actually accessing an i-node.
- file name look up is the same as we described before. The important thing is to find the root directory i-node, which is located in the fixed position of the disk.
- When user types in command  
*% mount /def/fd1 /user*  
this implies to mount the file system on top of /user.



# Directories

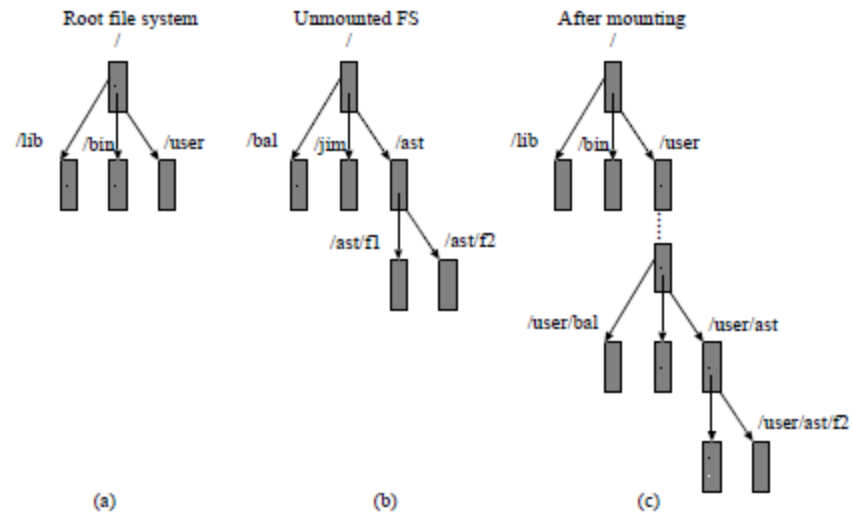


Figure 5.34: (a) Root file system. (b) An unmounted file system.  
(c) The result of mounting the file system of (b) on /user

# Directories

- What does it really mean for mounting? That is, during file name lookup, how does the system know when and where to switch?
- OS has to 1) set a flag in `"/user"` to indicate a successful mount, 2) load the super-block and i-node of the newly mounted file system into super-block table and i-node table (this is accomplished by MOUNT call), 3) set the *i-node-of-the-mounted-file-system* to point to the root i-node of the newly mounted system, 4) set the *i-node-mounted-on* to point to the i-node of `"/user"`.
- to answer the previous question, it is straight forward.

# Superblock

```
struct minix_super_block {  
    unsigned short s_ninodes;  
    unsigned short s_nzones;  
    unsigned short s_imap_blocks;  
    unsigned short s_zmap_blocks;  
    unsigned short s_firstdatazone;  
    unsigned short s_log_zone_size;  
    unsigned int s_max_size;  
    unsigned short s_magic;  
    unsigned short s_state;  
    unsigned int s_zones;  
};
```

# Inode

```
struct minix_inode {  
    unsigned short i_mode;  
    unsigned short i_uid;  
    unsigned int i_size;  
    unsigned int i_time;  
    unsigned char i_gid;  
    unsigned char i_nlinks;  
    unsigned short i_zone[9];  
};
```

# Directory Entry

```
struct minix_dir_entry {  
    unsigned short inode;  
    char name[0];  
};
```